

```
In [5]: import os  
import pandas as pd
```

```
In [7]: import seaborn as sns  
import matplotlib.pyplot as plt  
from scipy.stats import ttest_ind, skew, kurtosis, pearsonr, kruskal, f_oneway  
from statsmodels.stats.multicomp import pairwise_tukeyhsd  
from sklearn.preprocessing import StandardScaler  
from sklearn.cluster import KMeans  
from factor_analyzer import FactorAnalyzer  
from statsmodels.formula.api import ols  
import statsmodels.api as sm
```

```
In [9]: os.chdir('/Users/rosebui/desktop/Data science project/')
```

```
In [11]: df = pd.read_csv('HR_Data1.csv')
```

```
In [13]: df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3000 entries, 0 to 2999
Data columns (total 39 columns):
 #   Column           Non-Null Count Dtype  
--- 
 0   ID               3000 non-null   int64  
 1   FirstName        3000 non-null   object  
 2   LastName         3000 non-null   object  
 3   StartDate        3000 non-null   object  
 4   ExitDate         1533 non-null   object  
 5   Title            3000 non-null   object  
 6   Supervisor       3000 non-null   object  
 7   ADEmail          3000 non-null   object  
 8   BusinessUnit     3000 non-null   object  
 9   EmployeeStatus   3000 non-null   object  
 10  EmployeeType    3000 non-null   object  
 11  PayZone          3000 non-null   object  
 12  EmployeeClassificationType 3000 non-null   object  
 13  TerminationType 3000 non-null   object  
 14  TerminationDescription 1533 non-null   object  
 15  DepartmentType   3000 non-null   object  
 16  Division          3000 non-null   object  
 17  DOB              3000 non-null   object  
 18  State             3000 non-null   object  
 19  JobFunctionDescription 3000 non-null   object  
 20  GenderCode        3000 non-null   object  
 21  LocationCode      3000 non-null   int64  
 22  RaceDesc          3000 non-null   object  
 23  MaritalDesc       3000 non-null   object  
 24  Performance Score 3000 non-null   object  
 25  Current Employee Rating 3000 non-null   int64  
 26  Employee ID       3000 non-null   int64  
 27  Survey Date       3000 non-null   object  
 28  Engagement Score 3000 non-null   int64  
 29  Satisfaction Score 3000 non-null   int64  
 30  Work-Life Balance Score 3000 non-null   int64  
 31  Training Date     3000 non-null   object  
 32  Training Program Name 3000 non-null   object  
 33  Training Type     3000 non-null   object  
 34  Training Outcome   3000 non-null   object  
 35  Location           3000 non-null   object  
 36  Trainer            3000 non-null   object  
 37  Training Duration(Days) 3000 non-null   int64  
 38  Training Cost      3000 non-null   float64 
dtypes: float64(1), int64(8), object(30)
memory usage: 914.2+ KB

```

```
In [25]: date_cols = ["StartDate", "ExitDate", "DOB", "Survey Date", "Training Date"]
for col in date_cols:
    df[col] = pd.to_datetime(df[col], errors="coerce")
```

```
In [29]: df["ExitDate"].fillna(pd.NaT, inplace=True)
df.duplicated().sum()
```

```
/var/folders/xq/2xryw3gd7g7bx8y6s09b93780000gn/T/ipykernel_53975/1902752859.py:1: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method.  
The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.
```

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
df[\"ExitDate\"].fillna(pd.NaT, inplace=True)
```

```
Out[29]: 0
```

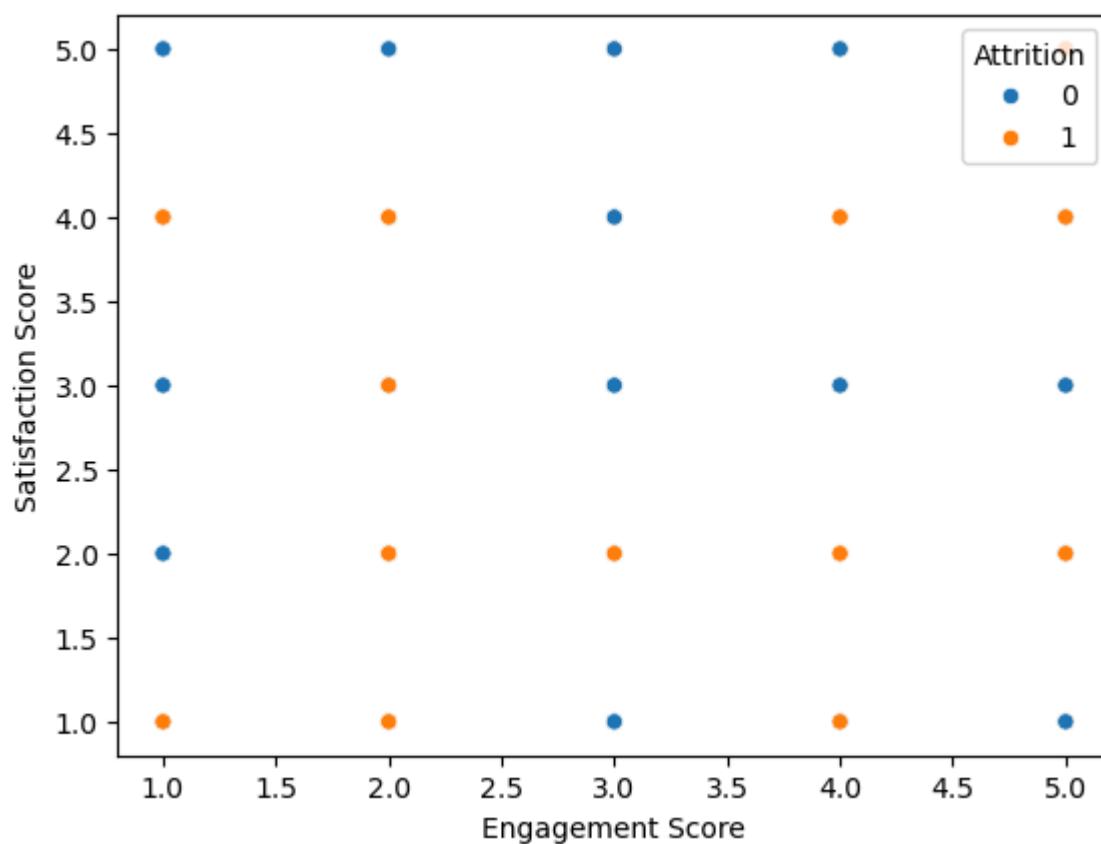
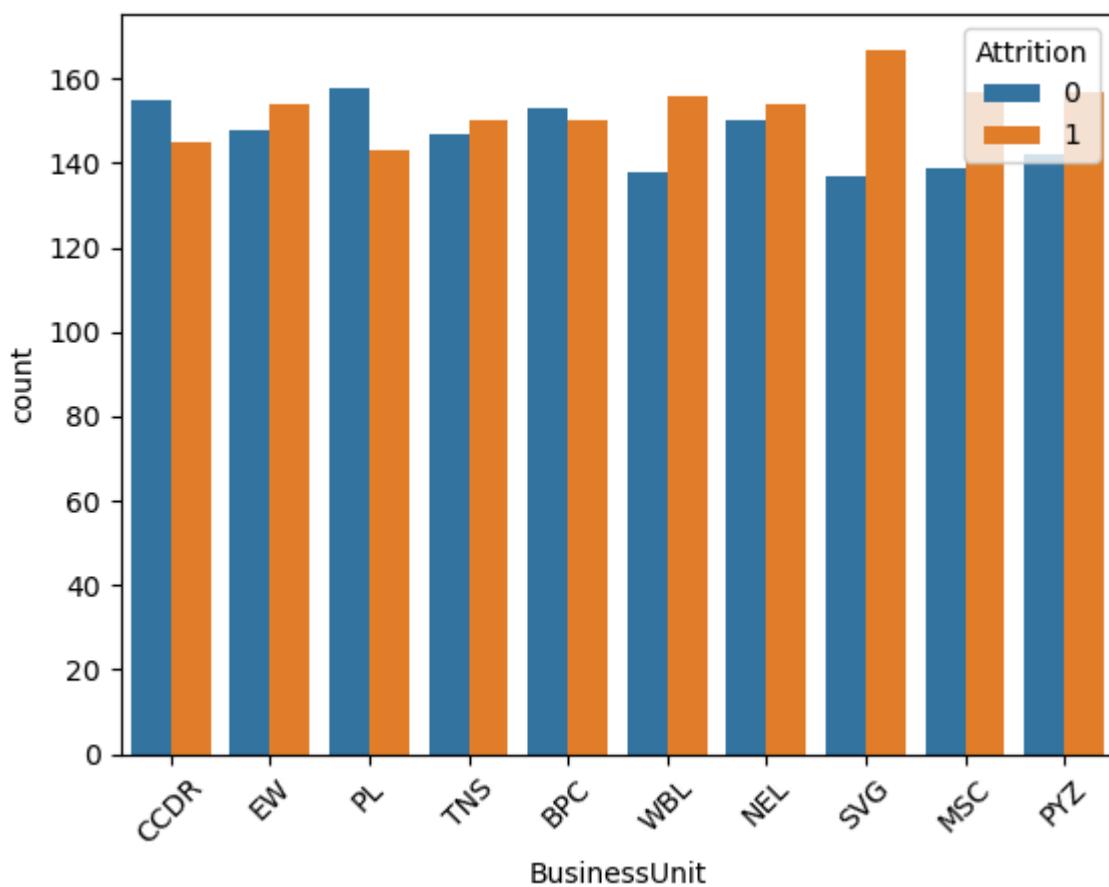
```
In [31]: df[\"TenureDays\"] = (df[\"ExitDate\"].fillna(pd.Timestamp.today()) - df[\"StartDate\"]).dt  
df[\"TenureYears\"] = df[\"TenureDays\"] / 365
```

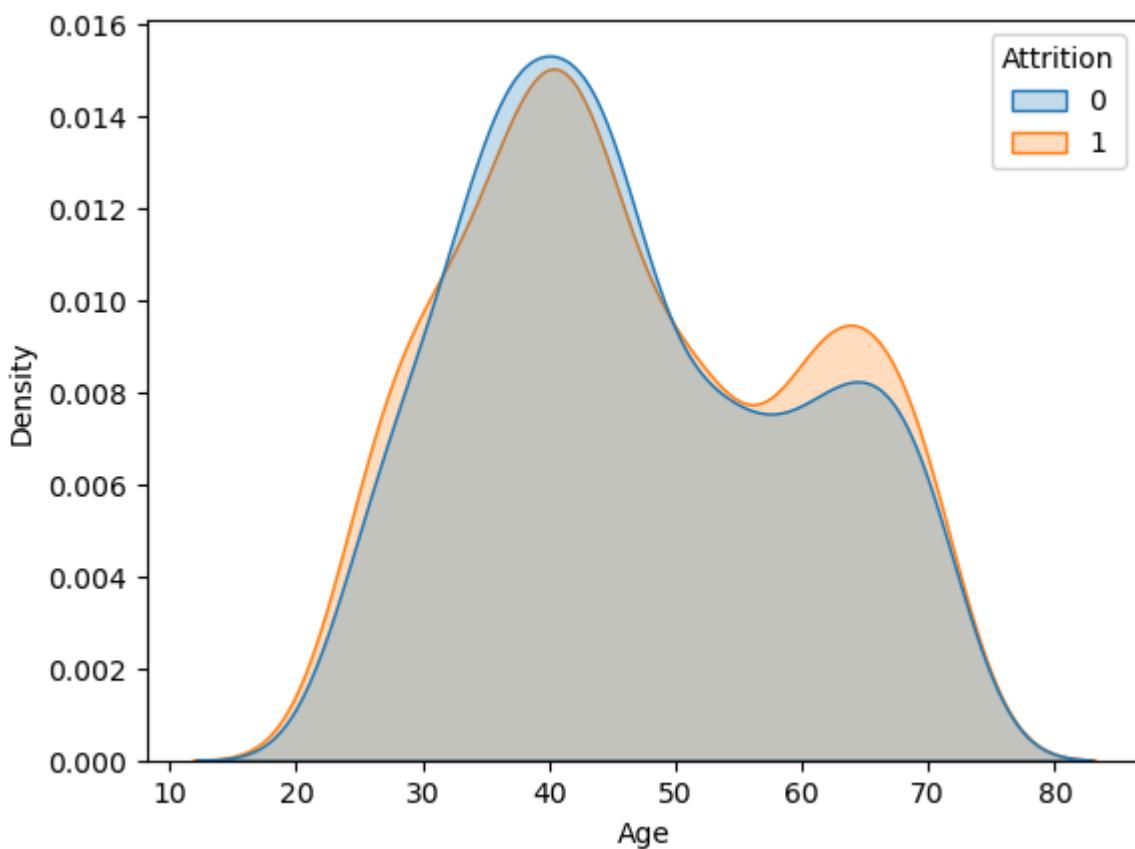
```
In [33]: df[\"Age\"] = (pd.Timestamp.today() - df[\"DOB\"]).dt.days // 365
```

```
In [35]: df[\"Attrition\"] = df[\"ExitDate\"].notnull().astype(int)
```

```
In [37]: df[\"TrainingCostPerDay\"] = df[\"Training Cost\"] / df[\"Training Duration(Days)\"]
```

```
In [39]: import seaborn as sns  
import matplotlib.pyplot as plt  
  
# Attrition by BusinessUnit  
sns.countplot(data=df, x="BusinessUnit", hue="Attrition")  
plt.xticks(rotation=45)  
plt.show()  
  
# Engagement vs Satisfaction  
sns.scatterplot(data=df, x="Engagement Score", y="Satisfaction Score", hue="Attrition")  
plt.show()  
  
# Age distribution by attrition  
sns.kdeplot(data=df, x="Age", hue="Attrition", fill=True)  
plt.show()
```





```
In [41]: from sklearn.model_selection import train_test_split
from sklearn.preprocessing import OneHotEncoder, StandardScaler
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import classification_report

# target
y = df["Attrition"]

# select features (example subset)
X = df[["Age", "TenureYears", "Engagement Score", "Satisfaction Score",
         "Work-Life Balance Score", "Performance Score", "GenderCode", "BusinessUnit"]]
```

```
In [43]: # categorical & numerical
cat_features = ["GenderCode", "BusinessUnit", "Performance Score"]
num_features = ["Age", "TenureYears", "Engagement Score", "Satisfaction Score", "Work

preprocessor = ColumnTransformer([
    ("cat", OneHotEncoder(handle_unknown="ignore"), cat_features),
    ("num", StandardScaler(), num_features)
])

# pipeline
model = Pipeline([
    ("preprocess", preprocessor),
    ("clf", RandomForestClassifier(random_state=42))
])

# train/test split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, stratify=y,

# fit
model.fit(X_train, y_train)

# evaluate
y_pred = model.predict(X_test)
print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.80	0.94	0.86	293
1	0.93	0.78	0.85	307
accuracy			0.85	600
macro avg	0.86	0.86	0.85	600
weighted avg	0.87	0.85	0.85	600

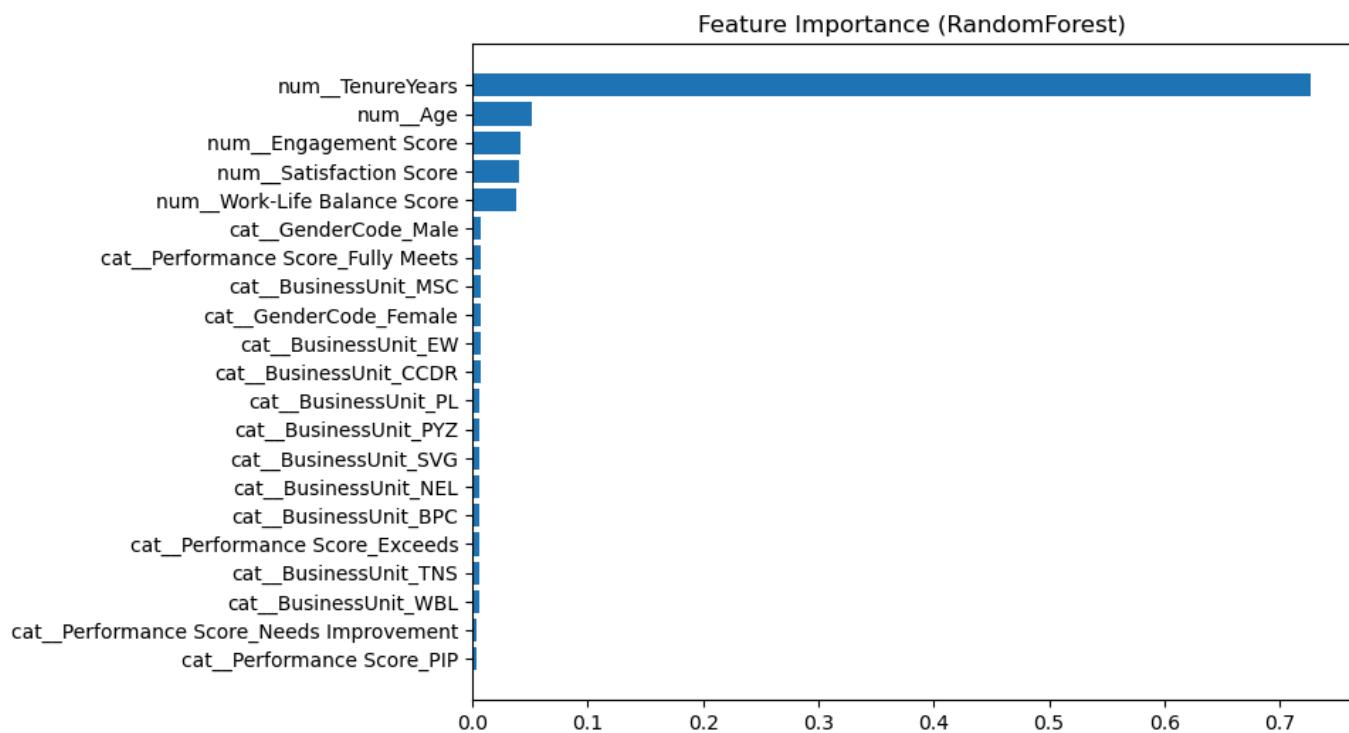
```
In [51]: import matplotlib.pyplot as plt
import pandas as pd

# Get feature names after preprocessing
feature_names = model.named_steps["preprocess"].get_feature_names_out()

# Get feature importance from RandomForest
importances = model.named_steps["clf"].feature_importances_

# Put into DataFrame
feat_imp = pd.DataFrame({
    "Feature": feature_names,
    "Importance": importances
}).sort_values("Importance", ascending=True)

# Plot
plt.figure(figsize=(8,6))
plt.barh(feat_imp["Feature"], feat_imp["Importance"])
plt.title("Feature Importance (RandomForest)")
plt.show()
```

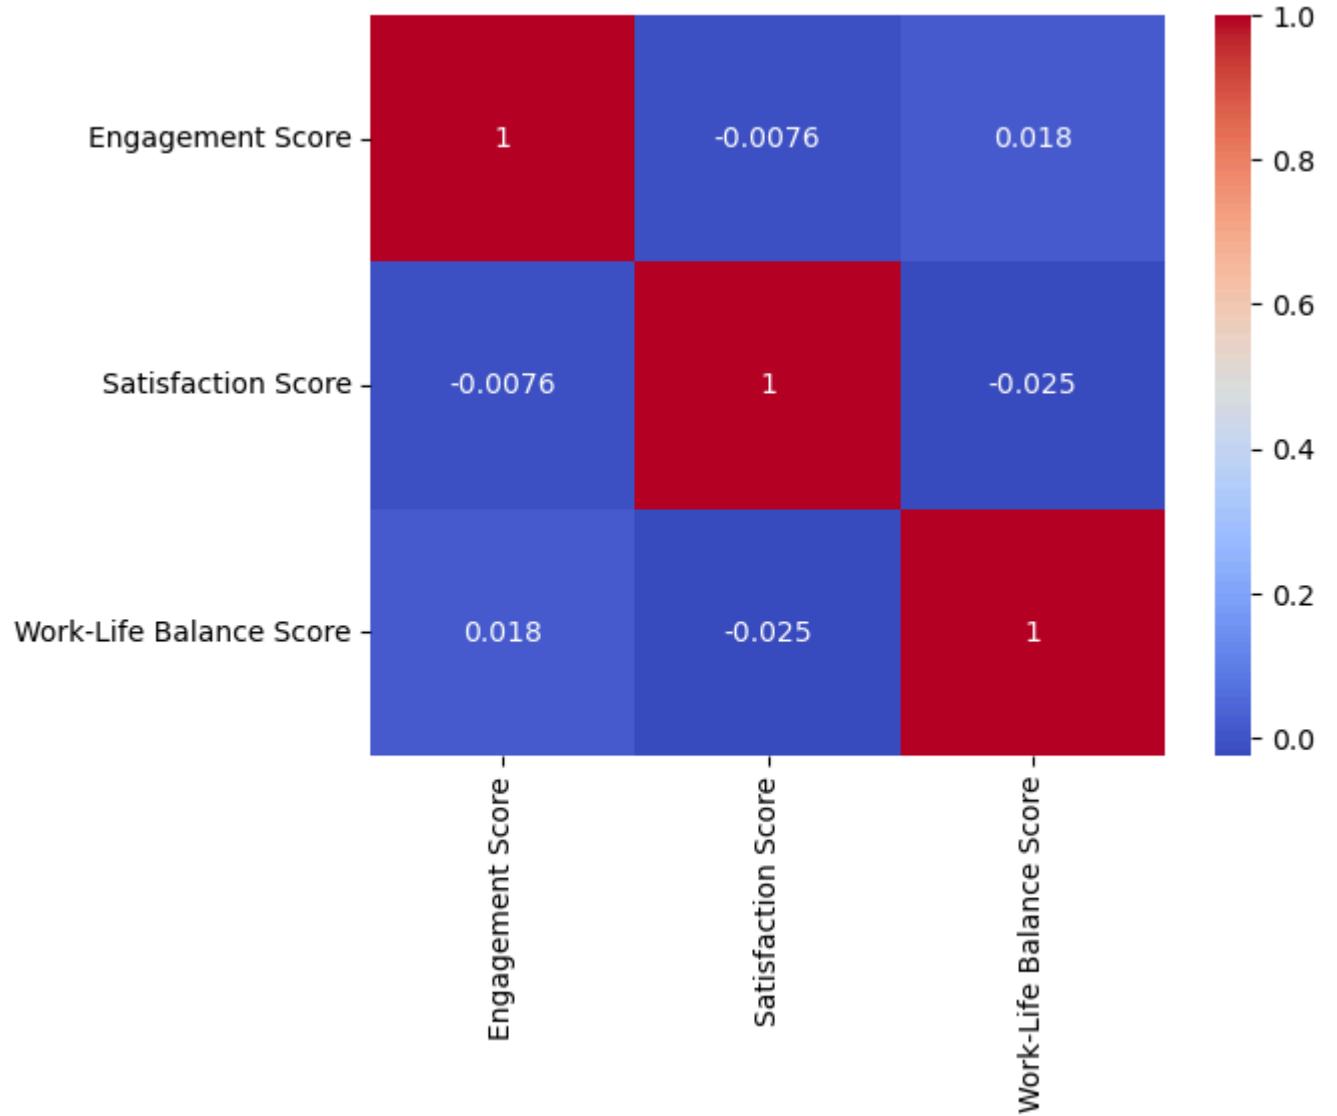


```
In [53]: # Attrition rate
attrition_rate = df["Attrition"].mean() * 100

# Attrition by department
dept_attrition = df.groupby("DepartmentType")["Attrition"].mean().sort_values() * 100
```

```
In [55]: import seaborn as sns
import matplotlib.pyplot as plt

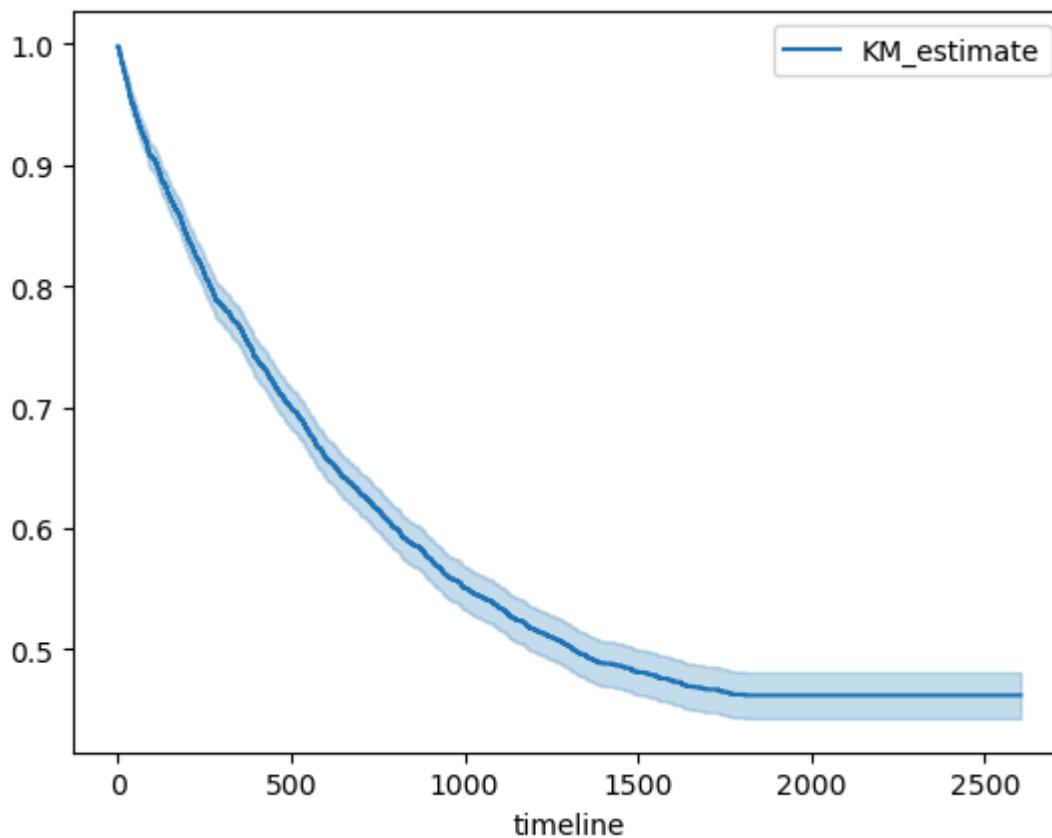
sns.heatmap(df[["Engagement Score", "Satisfaction Score", "Work-Life Balance Score"]])
plt.show()
```



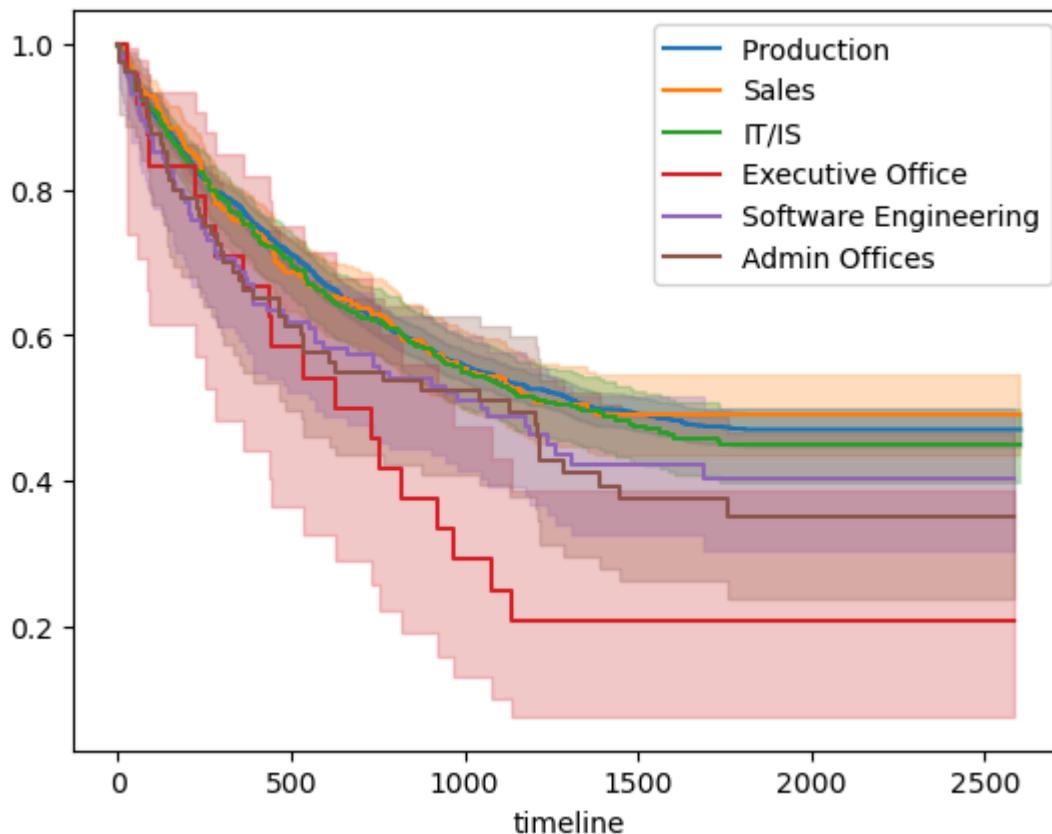
```
In [59]: from lifelines import KaplanMeierFitter
```

```
kmf = KaplanMeierFitter()  
kmf.fit(df["TenureDays"], event_observed=df["Attrition"])  
kmf.plot_survival_function()
```

```
Out[59]: <Axes: xlabel='timeline'>
```



```
In [83]: for dept in df["DepartmentType"].unique():
    kmf.fit(df[df["DepartmentType"]==dept]["TenureDays"],
            event_observed=df[df["DepartmentType"]==dept]["Attrition"],
            label=dept)
    kmf.plot_survival_function()
```



```
In [61]: from sklearn.cluster import KMeans
from sklearn.preprocessing import StandardScaler

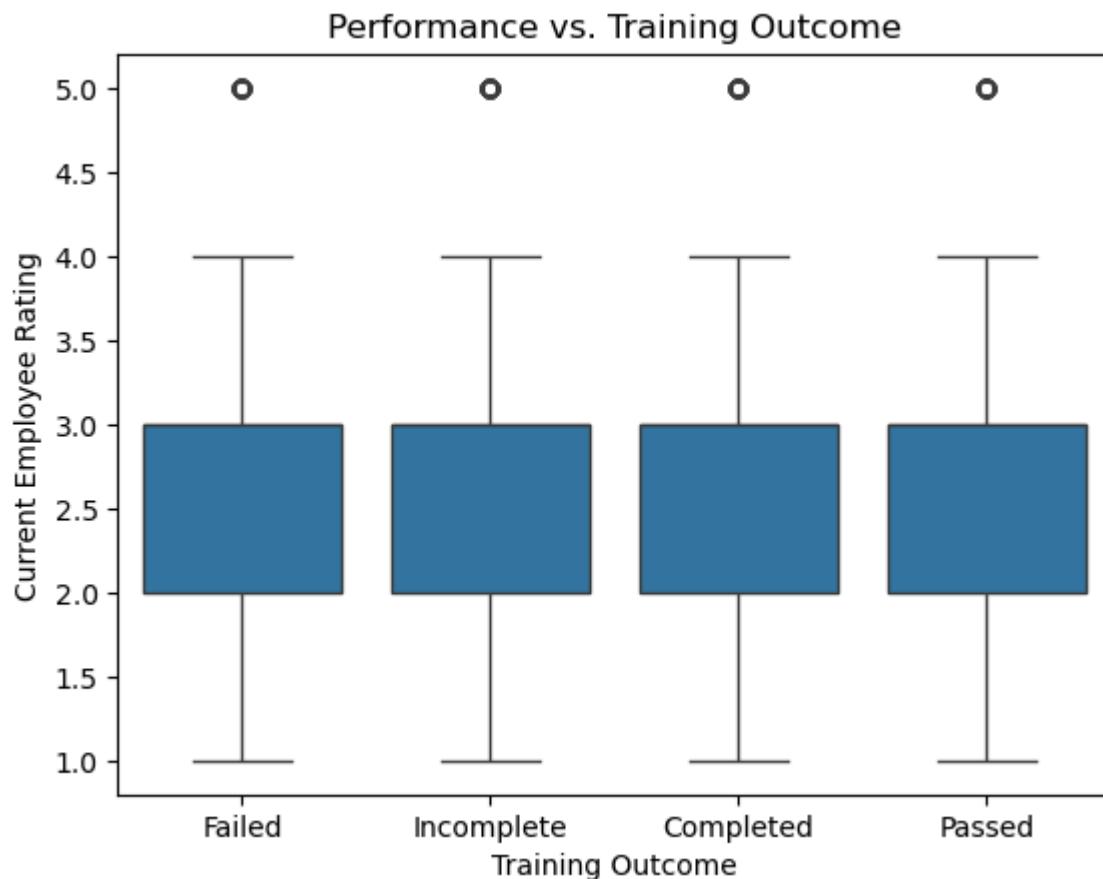
X = df[["Engagement Score", "Satisfaction Score", "Work-Life Balance Score"]]
X_scaled = StandardScaler().fit_transform(X)
```

```
kmeans = KMeans(n_clusters=3, random_state=42)
df[\"Cluster\"] = kmeans.fit_predict(X_scaled)
```

```
In [63]: import seaborn as sns
import matplotlib.pyplot as plt

sns.boxplot(data=df, x="Training Outcome", y="Current Employee Rating")
plt.title("Performance vs. Training Outcome")
plt.show()

# Average performance by training outcome
df.groupby("Training Outcome")["Current Employee Rating"].mean()
```



```
Out[63]: Training Outcome
Completed      3.003896
Failed        2.939944
Incomplete    2.989677
Passed        2.939107
Name: Current Employee Rating, dtype: float64
```

```
In [65]: sns.scatterplot(data=df, x="Training Cost", y="Engagement Score", hue="Training Outcome")
plt.title("Training Cost vs Engagement Score")
plt.show()

sns.scatterplot(data=df, x="Training Cost", y="Satisfaction Score", hue="Training Outcome")
plt.title("Training Cost vs Satisfaction Score")
plt.show()
```



Training ROI Analysis: This analysis evaluates whether training programs improve employee outcomes, including performance, satisfaction, and attrition. We also explore the relationship between training cost and ROI.

Step 1: Aggregate Training Outcomes: We calculate the average performance, satisfaction, attrition rate, and cost per training program. Then we create a simple **ROI Index** as:

$$\text{ROI Index} = (\text{Performance} + \text{Satisfaction}) - (\text{Attrition} \times 100)$$

```
In [67]: # Define ROI as performance gain + satisfaction gain - attrition rate
training_roi = df.groupby("Training Program Name").agg({
    "Current Employee Rating": "mean",
    "Satisfaction Score": "mean",
    "Attrition": "mean",
    "Training Cost": "mean"
})

# Create a simple ROI metric
training_roi["ROI Index"] = (training_roi["Current Employee Rating"] + training_roi["Satisfaction Score"] - training_roi["Attrition"]) / training_roi["Training Cost"]

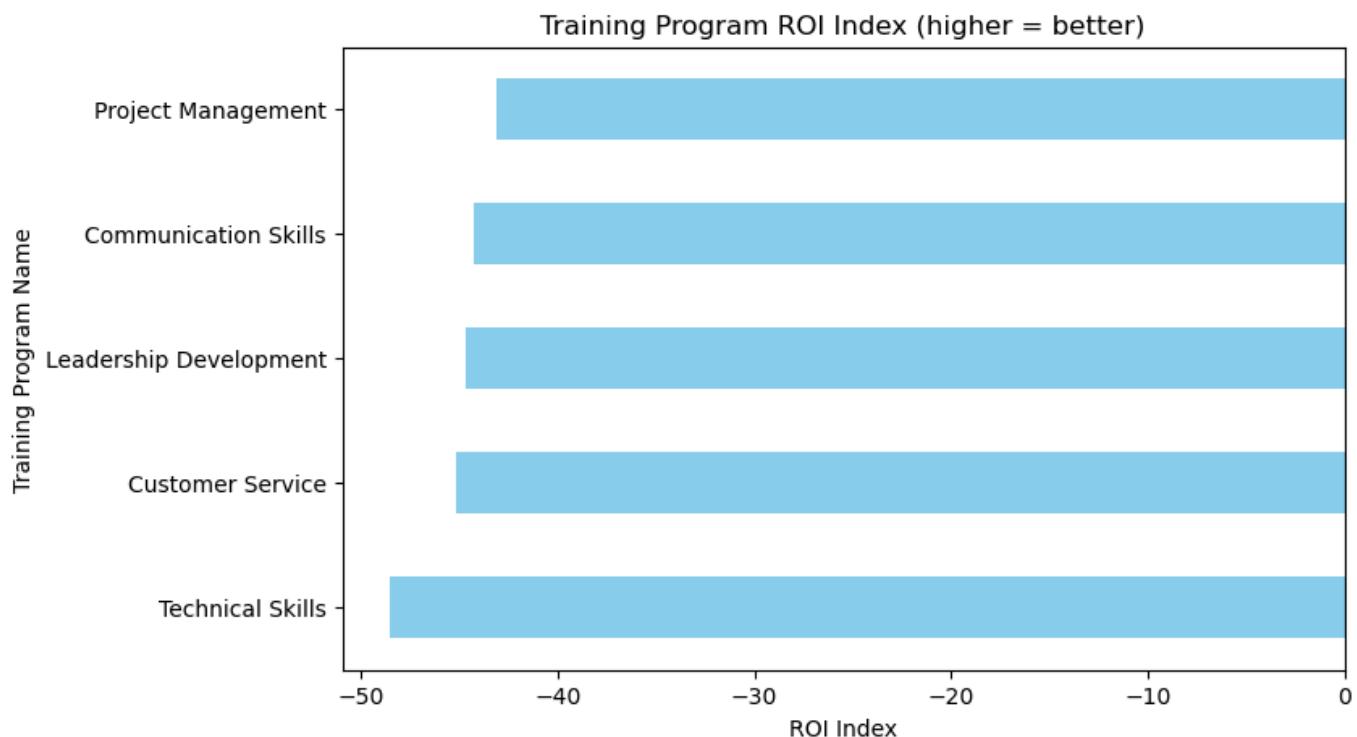
training_roi.sort_values("ROI Index", ascending=False)
```

Out[67]:

Training Program Name	Current Employee Rating	Satisfaction Score	Attrition	Training Cost	ROI Index
Project Management	2.954023	3.001642	0.490969	563.732627	-43.141215
Communication Skills	2.982169	2.968796	0.502229	542.382229	-44.271917
Leadership Development	2.949477	3.054007	0.506969	564.289251	-44.693380
Customer Service	2.969912	3.046018	0.511504	567.389451	-45.134513
Technical Skills	2.987910	3.050086	0.545769	557.983782	-48.538860

```
In [69]: import matplotlib.pyplot as plt
```

```
training_roi["ROI Index"].sort_values().plot(kind="barh", figsize=(8,5), color="skyblue")
plt.title("Training Program ROI Index (higher = better)")
plt.xlabel("ROI Index")
plt.show()
```

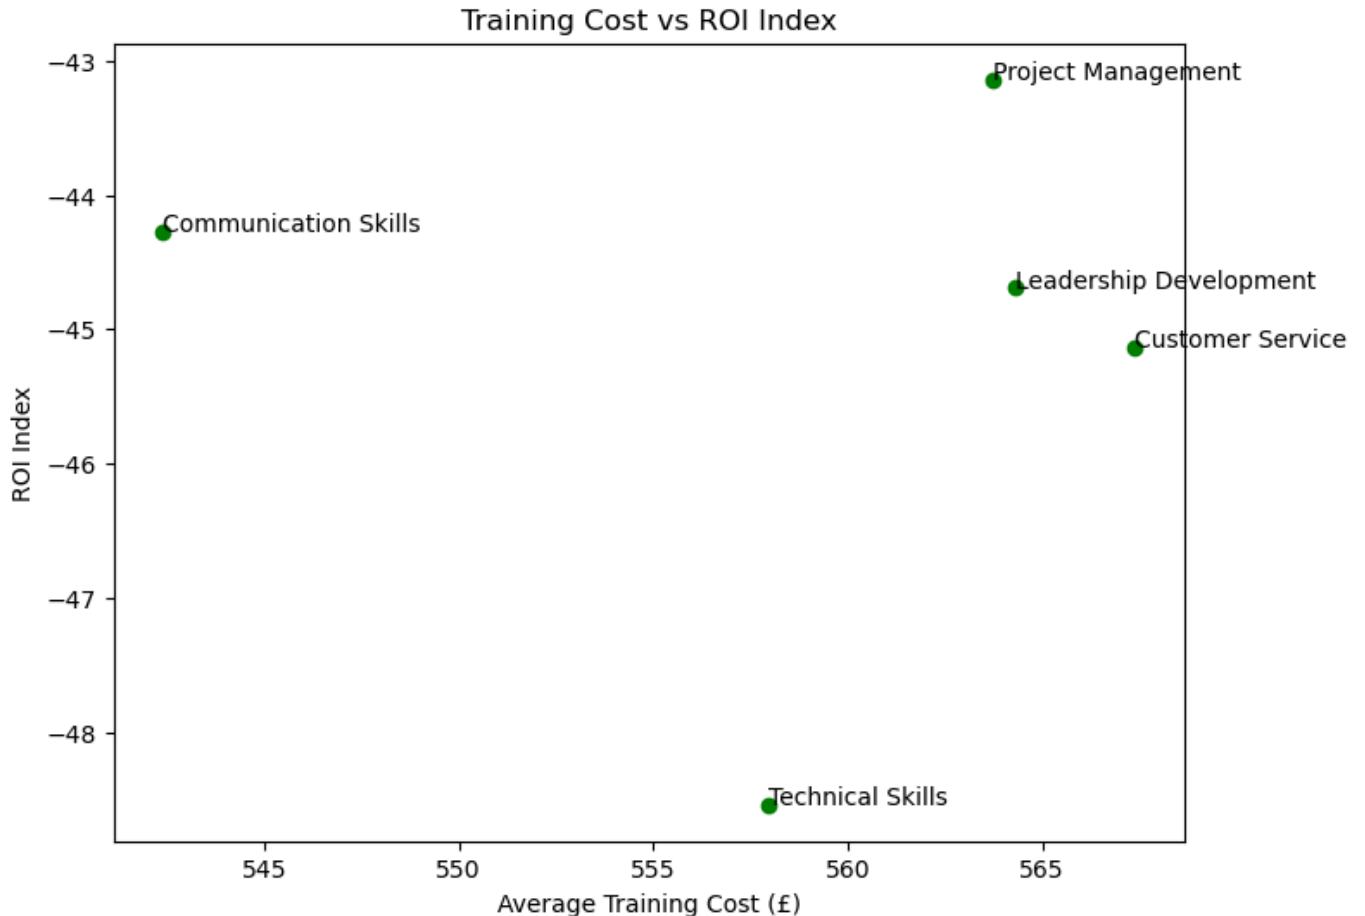


In [71]:

```
plt.figure(figsize=(8,6))
plt.scatter(training_roi["Training Cost"], training_roi["ROI Index"], color="green")
for i, txt in enumerate(training_roi.index):
    plt.annotate(txt, (training_roi["Training Cost"][i], training_roi["ROI Index"][i]))
plt.title("Training Cost vs ROI Index")
```

```
plt.xlabel("Average Training Cost (£)")  
plt.ylabel("ROI Index")  
plt.show()
```

```
/var/folders/xq/2xryw3gd7g7bx8y6s09b93780000gn/T/ipykernel_53975/1964518319.py:4: FutureWarning: Series.__getitem__ treating keys as positions is deprecated. In a future version, integer keys will always be treated as labels (consistent with DataFrame behavior). To access a value by position, use `ser.iloc[pos]`  
    plt.annotate(txt, (training_roi["Training Cost"][i], training_roi["ROI Index"][i]))
```



In []: